

Avaloka AI 1.0 — User Guide

- [Avaloka AI 1.0 — User Guide](#)
 - [About this guide](#)
 - [1. Welcome and core concepts](#)
 - [2. Getting started](#)
 - [3. Working with data](#)
 - [4. Visualization](#)
 - [5. Connections](#)
 - [6. Data Transfer](#)
 - [7. Model training and inference](#)
 - [8. Context Memory](#)
 - [9. Scheduling](#)
 - [10. FAQ and troubleshooting](#)
 - [Where to go next](#)

Avaloka AI 1.0 — User Guide

Document version 1.0 • Applies to Avaloka 1.6

About this guide

This is the **hands-on guide**: what you click, what you type, and what you see back. It walks through Avaloka the way you actually meet it — sign in, upload a file, ask a question, connect a database, train a model.

If you want the complete capability and integration reference instead — every format, every endpoint, every edition gate — read the [Technical User Guide](#). The two are companions: this one shows you the product, that one specifies it.

A note on screenshots. The images in this guide are placeholders. The previous guide's captures were taken against an earlier UI and no longer match what you will see, so they were removed rather than left to mislead — a stale screenshot is worse than none, because you trust it and then cannot find the control it shows. Each placeholder names the capture needed; see [images/user-guide/README.md](#).

1. Welcome and core concepts

Avaloka is an AI team of data scientists you direct in plain language. Give it a dataset and ask for what you want — “*find the fastest-growing channels*”, “*move this into my warehouse*”, “*train a model to predict churn*” — and Avaloka plans the analysis, writes and checks the code, runs it, and shows you the results.

You never write code, and nothing important happens without your say-so: you review plans and approve decisions along the way. One conversation can carry a dataset the whole distance — from a raw file, to cleaned and reshaped data, to charts and insights, to a trained model you can query for predictions.

What you can do

Analyze	Upload a dataset and ask questions in plain English. Each answer arrives as a new dataset you can keep refining.
Visualize	Charts appear automatically in the Auto Insights panel — chosen to fit your data, no configuration needed.
Move data	Register a database or cloud bucket as a connection, then ask Avaloka to transfer your dataset there, transforming it on the way.
Model and predict	Ask Avaloka to train a model, approve its plan, then test predictions in chat or from the Inference Panel.
Stay in context	Avaloka remembers your data’s structure, your preferences, and your past results, so each request builds on the last.

How it works

Behind the chat is a team of specialised agents led by a **Planner** — the agent you are actually talking to. When you send a prompt, the Planner decides what needs to happen and delegates: a **Coder** writes the Python, a **Validator** checks it against your data’s real structure and sends it back for fixes if anything is off, an **Execution** agent runs it, and a **Visualization** agent turns results into charts. Specialist agents handle transfers and model training, and heavy workloads run on a compute cluster rather than in your session.

You do not manage any of this. It is worth knowing only because this guide refers to these agents by name, and because it explains Avaloka’s rhythm: a moment of planning, then code, checks, execution, and results.

Your workspace

[Screenshot placeholder — dashboard-overview.png]

The Dashboard: your files on the left, data

and results in the middle, the AI Assistant alongside.

Area	What it is
Dashboard	Home base — files left, data and results centre, AI Assistant alongside
Upload File / File History	Add datasets and revisit earlier ones
Data Profile & Statistics	An automatic summary of each dataset: columns, types, missing values
AI Assistant	The chat panel where everything happens. Each file gets its own thread, listed under Chats
Auto Insights	Automatically generated charts for whatever data you are viewing
SQL Datastore / Cloud Dataset	Register connections to databases and cloud buckets
Inference Panel	Query your trained models for predictions

Words this guide uses

- **Dataset** — the analyzable data behind a file you uploaded or a source you connected. Every analysis produces a new one.
- **Prompt** — a plain-English request you type into the chat.
- **Plan** — Avaloka's proposed course of action, shown for review before big steps like model training.
- **Connection** — a saved link to a database or cloud bucket, referred to by a short alias you choose.
- **Transfer** — moving a dataset into a connection, with optional transformations on the way.
- **Model / run ID** — a trained predictor, and the receipt identifying the exact training run that produced it.
- **Session** — your current working context: your chat threads and what Avaloka currently remembers.

2. Getting started

Before you begin

You will need a modern browser — Chrome, Edge, Firefox, or Safari, kept reasonably up to date — and a valid work email address.

Self-hosting instead? See [INSTALL.md](#).

Creating your account

If your organisation already has a workspace, ask your administrator to invite you rather than signing up separately — that way your account is attached to the right team and its data.

1. Go to the Avaloka sign-in page.
2. Click **Sign Up / Get Started** in the top-right corner.
3. Enter your **Email**.
4. Enter a **Password**.
5. Re-enter it in **Confirm Password**.
6. Enter your full name and organisation.
7. Review the Terms of Service and Privacy Policy, then accept them.
8. Click **Create Account**.

Avaloka sends a verification email. **You cannot sign in until you verify your address.** Open the message and click **Verify Email**. If it has not arrived within a few minutes, check your spam folder, or click **Resend email** on the verification screen.

Logging in

1. Click **Log In** in the top-right corner.
2. Enter your registered email and password.
3. *(Optional)* Select **Remember me** to stay signed in on this device. Leave it clear on shared or public computers.
4. Click **Log In**.

First time in, you will see an onboarding flow and an empty workspace.

Resetting a forgotten password

1. On the login page, click **Forgot password?**
2. Enter your registered email and click **Send reset link**.
3. Open the reset email and click **Reset Password**. The link is single-use and expires.
4. Enter and confirm your new password, then click **Update Password**.

For security, Avaloka shows the same confirmation whether or not the address is registered. If no email arrives, the likeliest cause is that you signed up with a different address.

3. Working with data

Uploading a file

Click **Upload File**. Accepted formats are **CSV, Excel, JSON, and Parquet**, up to **100 MB per file**.

[Screenshot placeholder — `upload-complete.png`] *A finished upload with the dataset open.*

The left panel lists everything you have uploaded; clicking a file opens the dataset.

[Screenshot placeholder — `file-history.png`] *File History listing previously uploaded datasets.*

Reading the Data Profile

The moment a dataset arrives, Avaloka gets to know it — counting and typing every column, measuring what is missing, and explaining what each column means. The result appears before you have asked a single question.

Every upload runs the same pipeline automatically:

1. **Upload** — the file lands in Avaloka's storage.
2. **Sampling** — a portfolio of representative samples is built so browsing and reasoning stay fast.
3. **Profiling** — every column is measured: type, missing values, distinct values, ranges.
4. **AI analysis** — column meanings and readiness are explained.

[Screenshot placeholder — `data-profile-card.png`] *The profile card above a dataset.*

Click **Show More** for the full briefing:

- **Score and reasoning** — why the readiness score is what it is
- **Key Relationships** — pairs of columns expected to move together
- **Red Flags** — problems to know about before you trust your results
- **Quick Wins** — small clean-up steps worth doing first

[Screenshot placeholder — `quick-insights-expanded.png`] *Quick Insights expanded: score, relationships, red flags, quick wins.*

Prompting

Instead of building formulas, ask for what you want. Avaloka creates a plan based on your prompt and decides which agents are needed — ask it to train a model and the Model Training agent is routed in; ask it to reshape a dataset and the Coder and Validator handle it.

Try something concrete:

“Identify the fastest-growing channels by comparing `subscribers_for_last_30_days` to total subscribers. List the top 20 and include Youtuber, category, Country, subscribers, and `subscribers_for_last_30_days`.”

When Avaloka finishes, the result appears in the output view as a **new dataset**. Your original is untouched — every analysis produces a new dataset you can keep refining.

[Screenshot placeholder — prompt-and-output.png] *A prompt and the output dataset it produced.*

If a prompt produces no dataset — a question answered in prose, for example — you will see the answer in the chat area instead.

[Screenshot placeholder — no-output-message.png] *The chat message shown when a prompt produces no output dataset.*

Sampled or complete?

For speed, Avaloka explores on a sample. When you need an exact figure, say so:

“How many distinct aisles are there?” — *fast, sampled* “Now compute that on the entire dataset.” — *exact, full scan*

Avaloka always tells you which it used. This matters: a sampled distinct-count is an estimate, and treating it as exact is how a confident wrong answer happens.

4. Visualization

Avaloka creates charts for you automatically. Whenever you upload a dataset or run an analysis, the **Auto Insights** panel fills with **three to five charts** chosen to explain the data in front of you. You do not have to ask, or pick chart types.

[Screenshot placeholder — auto-insights-panel.png] *Auto Insights with automatically generated charts.*

Every prompt that produces a result dataset also refreshes the charts. Asking a chart-shaped question steers what gets visualized:

“Plot monthly revenue by region.” “Chart trips by hour of day.” “Scatter fare vs distance.” “Compare holiday vs non-holiday sales visually.”

Charts are interactive — hover for values, click the expand icon for fullscreen.

[Screenshot placeholder — chart-expanded.png] *A chart opened fullscreen.*

Chart	Best for
Bar	Comparing categories, such as sales per region
Line	Change over time or another ordered axis
Pie	How a whole splits into parts, for a small number of categories

5. Connections

A connection stores the address and sign-in details for one place your data lives — a database or a cloud storage bucket. You create it once and refer to it afterwards by the name you gave it. Connections are created in the **Data Sources** panel.

Connections to external databases and clouds are a **paid capability**. File upload and analysis are available in every edition. See [EDITIONS.md](#).

Connecting to a database

In **Data Sources**, expand **SQL Datastore** → **Connections** → **New Connection**.

[Screenshot placeholder — `sql-connection-form.png`] *The Database Connection Details form.*

Field	What to enter
Connection Name *	A name for this connection, e.g. <code>Production Database</code>
Customer ID *	A short identifier — letters, numbers and underscores only
Customer Name *	Your organisation's full name
Contact Email *	A contact address for this connection
Connection Type *	The purpose of this database, e.g. <code>production</code>
Database Type *	Select from the dropdown
Host *	The server address, e.g. <code>db.example.com</code>
Port	The port number, e.g. <code>5432</code>
Database Name *	The database to connect to
Username *	The sign-in username
Password *	The sign-in password
Query Timeout (seconds)	How long a query may run before stopping
Max Rows	The maximum rows a query returns

Fields marked * are required. Click **Test & Register** — Avaloka tests the connection before saving, so an incorrect detail is reported immediately rather than during a later transfer.

[Screenshot placeholder — connection-dashboard.png] *The Connection Dashboard after a successful register.*

Connecting to cloud storage

Under the **Cloud** tab, click **New Connection**. Avaloka supports **AWS, Azure, and GCP**.

[Screenshot placeholder — cloud-connection-form.png] *The cloud connection modal.*

After **Test & Save Connection**, your connections are listed with actions: the folder icon browses datasets inside the connection, the pen icon updates the details.

[Screenshot placeholder — cloud-browser.png] *Browsing datasets inside a cloud connection.*

6. Data Transfer

The **Data Transfer Agent (DTA)** moves data from one place to another — into a database, into a cloud bucket, or into a different file format — and can clean or reshape it on the way. You describe the transfer in chat; the agent writes and runs it.

“Transfer aisles.csv from the instacart-mba connection to the walmart connection under transfers/aisles_copy.csv.” “Move my cleaned housing dataset to the nyc-taxi bucket connection.”

[Screenshot placeholder — dta-transfer-chat.png] *A transfer requested and confirmed in chat.*

What it can read

Source	Supported
Files	CSV, JSON, Parquet
Databases	MySQL, PostgreSQL

CSV and Parquet can be read **directly from a cloud bucket**. JSON can only be read after it has been loaded into Avaloka — if your JSON is in a bucket, convert it to CSV or Parquet first.

What it can write

Destination	Supported
Files	CSV, TSV, JSON, Parquet, Excel (.xlsx), XML, Avro, ORC
Databases	MySQL, PostgreSQL

All eight file formats are available when writing to a cloud bucket.

Read and write are not symmetric. Avaloka writes eight file formats but reads three. If you need to transfer *from* a format not in the read list, load it into Avaloka first, then transfer.

Databases you can query but not transfer into

You can register and query **SQLite, SQL Server, Oracle, and MariaDB**. Transfers currently target **MySQL and PostgreSQL** only.

7. Model training and inference

The **Model Training Agent (MTA)** turns a dataset into a model you can test in chat. Tell it what to predict, review its choices, approve training, then ask for predictions.

Worked example

Dataset: `housing.csv` (California Housing Prices) **Target:** `median_house_value` · **Task:** regression

1. Upload the dataset. Confirm `median_house_value` appears as a column before requesting training.

2. Ask for training in one direct request:

“Train a regression model to predict `median_house_value` using longitude, latitude, `housing_median_age`, `total_rooms`, `total_bedrooms`, population, households, `median_income`, and `ocean_proximity`.”

3. Review the plan. It shows the model type, target, selected features, epochs, optimizer, learning rate, and batch size. **Training has not started yet.**

[Screenshot placeholder — `mta-training-plan.png`] *The training plan shown for review.*

4. Change or expand it. Ask “Show the full detailed training plan” for everything, or amend it conversationally — “Add `SibSp` as a feature too”, “Change it to 5 epochs”.

5. Approve:

“Looks good — go ahead and train it.”

[Screenshot placeholder — `mta-training-progress.png`] *Training in progress.*

6. Predict. Provide example values in chat, or use the **Inference Panel**.

[Screenshot placeholder — `inference-panel.png`] *The Inference Panel with a prediction returned.*

Before training begins, Avaloka screens the data for **leakage** — a feature that gives away the answer. If it finds any, training is blocked and the finding explained. A leaking model reports excellent scores and fails in production, so this is a refusal worth having.

Every result is compared against a **trivial baseline**. If a model does not beat it, Avaloka says so plainly — that is a real finding, not a failure.

8. Context Memory

Context Memory lets Avaloka remember your data, preferences, and past results while you work, so each prompt builds on what came before. It works automatically, and you can also tell it explicitly what to remember.

What Avaloka remembers:

- **Your dataset** — its columns and layout, so you do not re-describe it every prompt.
- **Your working style** — built up quietly as you work; for example, that you prefer medians, or always exclude null rows.
- **Things you ask it to remember** — any preference or fact you state explicitly. These are **protected**: newer memories never crowd them out.
- **Your past results** — charts, tables, and model results it can reuse instead of redoing the work.

“Preference: always include row counts in any table you show me.” “Use the cleaned dataset from before.” “What did we find about fraud rates in earlier analyses?”

9. Scheduling

Append *when to run* to any prompt and Avaloka creates a scheduled job:

“Every Monday 9am, recompute the fraud rate by ProductCD and store the result.”

[Screenshot placeholder — scheduler-task-created.png] *A scheduled job created from a prompt.*

For a repeating task, say how often to run it and how many times to repeat. Any phrasing of the schedule works.

Ask for your tasks in chat at any time:

“List my scheduled tasks.”

[Screenshot placeholder — scheduler-task-list.png] *The task list returned in chat.*

When a job finishes, its status message updates. If the job modified a dataset, the output view refreshes automatically.

Model training, data transfers, distributed execution, and inference-service startup all schedule their own jobs when requested through chat.

10. FAQ and troubleshooting

What you see	What to do
Cannot sign in after registering	Verify your email first — the account is inactive until you do
No verification or reset email	Check spam; the likeliest cause is a different address at signup
A capability is refused	It may be gated by your edition — see EDITIONS.md
Two answers give different numbers	One was sampled, one was complete. Ask “on the entire dataset” for the authoritative figure
Training will not start	The integrity screen found leakage. Read the finding — a leaking feature usually should be dropped
A model reports poor results	It did not beat the baseline. The signal may not be in the data
A transfer will not read my file	Check the read formats in §6 — reads are narrower than writes
Upload rejected	Files are limited to 100 MB; use a connection for larger data

Always log out on shared devices: click your profile icon in the top-right and select **Log Out**.

Where to go next

Every capability, format, endpoint and edition gate	Technical User Guide
Installing and licensing	INSTALL.md · EDITIONS.md
Calling Avaloka from code	api.md · cli.md
Running it on Kubernetes	deployment.md

Avaloka is open source under the Apache License 2.0.