

Grounding — the plan is built from the data, not from the chat log

Analyst question
chat · REST
CLI · MCP

Conversational
agent +
episodic memory

Sampling +
profiling
schema · stats

Planner +
Summarizer →
planner graph

Code loop — execution sits *between* the validators

Coder
blueprint +
contract

1 syntax
2 static AST
20 checks

execute on a
100-row sample

3 contract check
blocks
4 LLM review
advisory

refine

Evidence gates — the answer is checked, not only the code

Execute
quick_sample → local
portfolio_samples → Daft/Ray
entire_dataset → Ray on k8s

Integrity
leakage check
before training

Claim verifier
narrative vs
evidence
no LLM

Answer +
artifacts +
verification
report

Model routing — one chokepoint, shared by all 12 agents above

model_config
→ which model
build_chat_model
→ which provider

groq · local vLLM
openai · openrouter
bedrock · vertex
azure